

[B I T S A N D B Y T E S] B i t s a n d b y t e

Bitsandbytes documentation bitsandbytes and get a
experience to get started bitsandbytes enables a
quantization for PyTorch

来源：<https://huggingface.co/docs/bitsandbytes/index>

分类：12_模型部署与工程化

Bitsandbytes documentation

bitsandbytes

and get access to the augmented documentation exp

to get started

bitsandbytes enables accessible large language mo
bitsandbytes provides three main features for dra
inference and training:

8-bit optimizers uses block-wise quantization to
fraction of the memory cost.

LLM.int8() or 8-bit quantization enables large la
required memory and without any performance degra
quantization to quantize most features to 8-bits
16-bit matrix multiplication.

QLoRA or 4-bit quantization enables large languag
techniques that don't compromise performance. Thi
inserts a small set of trainable low-rank adaptat

bitsandbytes is MIT licensed.