


Join the Hugging Face community

and get access to the augmented documentation experience

Sign Up

to get started



D[🧨]ffusers

Diffusers

Copy page



Diffusers is a library of state-of-the-art pretrained diffusion models for generating videos, images, and audio.

The library revolves around the [DiffusionPipeline](#), an API designed for:

- easy inference with only a few lines of code
- flexibility to mix-and-match pipeline components (models, schedulers)
- loading and using adapters like LoRA

Diffusers also comes with optimizations - such as offloading and quantization - to ensure even the largest models are accessible on memory-constrained devices. If memory is not an issue, Diffusers supports `torch.compile` to boost inference speed.

Get started right away with a Diffusers model on the [Hub](#) today!

Learn

If you're a beginner, we recommend starting with the [Hugging Face Diffusion Models Course](#). You'll learn the theory behind diffusion models, and learn how to use the Diffusers library to generate images, fine-tune your own models, and more.

[Update](#) on GitHub

Installation →