

[L I G H T E V A L] L i g h t e v a l

L i g h t e v a l d o c u m e n t a t i o n L i g h t e v a l a n d g e t a c c e s s
t o g e t s t a r t e d L i g h t e v a l i s y o u r a l l - i n - o n e t o
(L L M s) a c r o s s m u l t . . .

来源：<https://huggingface.co/docs/lighteval/index>

分类：11_数据处理与工程

L i g h t e v a l d o c u m e n t a t i o n

L i g h t e v a l

a n d g e t a c c e s s t o t h e a u g m e n t e d d o c u m e n t a t i o n e x p
t o g e t s t a r t e d

L i g h t e v a l i s y o u r a l l - i n - o n e t o o l k i t f o r e v a l u
(L L M s) a c r o s s m u l t i p l e b a c k e n d s w i t h e a s e . D i v e o
p e r f o r m a n c e b y s a v i n g a n d e x p l o r i n g d e t a i l e d , s a m
a n d s e e h o w y o u r m o d e l s s t a c k u p .

K e y F e a t u r e s

M u l t i - B a c k e n d S u p p o r t

E v a l u a t e y o u r m o d e l s u s i n g t h e m o s t p o p u l a r a n d e

e v a l : U s e i n s p e c t - a i a s b a c k e n d t o e v a l u a t e a n d i
e v a l

t r a n s f o r m e r s : E v a l u a t e m o d e l s o n C P U o r o n e o r m o
A c c e l e r a t e

t r a n s f o r m e r s

n a n o t r o n : E v a l u a t e m o d e l s i n d i s t r i b u t e d s e t t i n g s
N a n o t r o n

nanotron

vllm: Evaluate models on one or more GPUs using VLLM

vllm

custom: Evaluate custom models (can be anything)

custom

sglang: Evaluate models using SGLang as backend

sglang

inference-endpoint: Evaluate models using Hugging

inference-endpoint

tgi: Evaluate models using Text Generation Inf

tgi

litellm: Evaluate models on any compatible API us

litellm

inference-providers: Evaluate models using Huggin
Distributed training and evaluation

inference-providers

Comprehensive Evaluation

Extensive Task Library: 1000s pre-built evaluation

Custom Task Creation: Build your own evaluation t

Flexible Metrics: Support for custom metrics and

Detailed Analysis: Sample-by-sample results for c

Easy Customization

Customization at your fingertips: create new task metrics or model tailored to your needs, or browse

Seamless Integration

Seamlessly experiment, benchmark, and store your locally.

Quick Start

Installation

```
pip install lighteval
```

Basic Usage

```
lighteval eval "hf-inference-providers/openai/gpt-4o-mini:  
gpqa:diamond" \  
--bundle-dir gpt-oss-bundle \  
--repo-id OpenEvals/evals
```

Resulting Space: